

## CSE 512 Spring 2021 – Machine Learning – Homework 6

Name: Irfan Ahmed

Solar ID: 113166464

NetID Email: [irfan.ahmed@stonybrook.edu](mailto:irfan.ahmed@stonybrook.edu)

1.1)a)

```
X = tensor([[0.3868, 0.7403],
            [0.6417, 0.5854],
            [0.1191, 0.1833]], dtype=torch.float64)
```

```
b) Y = tensor([[1., 1.],
              [1., 1.],
              [1., 1.]], dtype=torch.float64) torch.Size([3, 2])
```

```
c) out = tensor([[1.3868, 1.7403],
                 [1.6417, 1.5854],
                 [1.1191, 1.1833]], dtype=torch.float64)
```

In-place addition :

```
Y = tensor([[1.3868, 1.7403],
            [1.6417, 1.5854],
            [1.1191, 1.1833]], dtype=torch.float64)
```

d) Randomly generated numpy array :

```
X = [[0.35990055 0.82630721]
      [0.25775084 0.98181074]
      [0.46863991 0.85011467]]
```

Converted to tensor :

```
tensor([[0.3599, 0.8263],
        [0.2578, 0.9818],
        [0.4686, 0.8501]], dtype=torch.float64)
```

Converted back to numpy array:

```
[[0.35990055 0.82630721]
 [0.25775084 0.98181074]
 [0.46863991 0.85011467]]
```

1.2) Requires\_grad = true:

```
tensor([[0.0891, 0.2565],
        [0.9623, 0.7975],
        [0.9608, 0.4749]], requires_grad=True)
```

b) Multiplied by 10:

```
tensor([[0.8914, 2.5654],
        [9.6231, 7.9753],
        [9.6077, 4.7492]], grad_fn=<MulBackward0>)
```

Added 0.1:

```
tensor([[0.9914, 2.6654],
        [9.7231, 8.0753],
        [9.7077, 4.8492]], grad_fn=<AddBackward0>)
```

Taking max:

```
tensor(9.7231, grad_fn=<MaxBackward1>)
```

c) Gradient:

```
tensor([[ 0.,  0.],
        [10.,  0.],
        [ 0.,  0.]])
```

d) After multiply:

```
tensor([[0.8914, 2.5654],
        [9.6231, 7.9753],
        [9.6077, 4.7492]])
```

After addition:

```
tensor([[0.9914, 2.6654],
        [9.7231, 8.0753],
        [9.7077, 4.8492]])
```

Max:

```
tensor(9.7231)
```

Since no grad(), these cannot be traced back to do backpropagation.

### 1.3) a)

Model:

```
TNet(  
  (conv1): Conv2d(3, 64, kernel_size=(3, 3), stride=(1, 1))  
  (layer1): Sequential(  
    (0): ReLU()  
    (1): MaxPool2d(kernel_size=4, stride=4, padding=0, dilation=1,  
ceiling_mode=False)  
  )  
  (conv2): Conv2d(64, 128, kernel_size=(3, 3), stride=(1, 1))  
  (layer2): Sequential(  
    (0): Sigmoid()  
    (1): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,  
ceiling_mode=False)  
  )  
  (classifier): Sequential(  
    (0): Linear(in_features=512, out_features=10, bias=True)  
  )  
  (dropout): Dropout(p=0.5, inplace=False)  
)
```

No. of parameters: 6

### b) Output and size:

```
tensor([[ -0.2261,  0.3337, -0.1659, -0.1669, -0.4315,  0.0406, -0.2984,  
  0.0986, -0.3152,  0.0104]], grad_fn=<AddmmBackward>) torch.Size([1, 10])
```

### c)

Gradients of intermediate layer:

layer 2.bias.grad after backward

```
tensor([-0.0018, 0.0000, 0.0269, 0.0092, 0.0123, 0.0231, 0.0108, -0.0168, 0.0000,  
  0.0000,  0.0080, -0.0099, -0.0008,  0.0065,  0.0089,  0.0280, -0.0008,  
  0.0023,  0.0046, -0.0200, -0.0044, -0.0173,  0.0164, -0.0146, 0.0158, -  
  0.0067,  0.0062,  0.0109, -0.0244,  0.0067,  0.0000, -0.0122, 0.0011,  
  0.0000, -0.0169, -0.0143,  0.0087,  0.0038,  0.0048,  0.0074, -0.0180, -  
  0.0006, -0.0098,  0.0000,  0.0000,  0.0025,  0.0065, -0.0015,  
  0.0064,  0.0112,  0.0213, -0.0004,  0.0027,  0.0147,  0.0146, -0.0004,  
 -0.0121,  0.0274, -0.0044,  0.0045,  0.0000,  0.0007, -0.0164,  0.0093,  
  0.0000,  0.0110, -0.0063, -0.0286, -0.0018,  0.0142, -0.0017, -0.0054,  
 -0.0013,  0.0041,  0.0275,  0.0161,  0.0000, -0.0291,  0.0017,  0.0107,  
  0.0000,  0.0098,  0.0006, -0.0077, -0.0089,  0.0022,  0.0000,  0.0113,  
  0.0135,  0.0082,  0.0109, -0.0034,  0.0082, -0.0023,  0.0000, -0.0034,  
  0.0105, -0.0014,  0.0001, -0.0102,  0.0129,  0.0035, -0.0030, -0.0134,  
 -0.0071, -0.0041,  0.0000,  0.0022, -0.0043, -0.0311,  0.0045,  0.0000,  
  0.0142, -0.0071, -0.0028, -0.0153,  0.0127, -0.0080,  0.0080, -0.0109,  
  0.0001,  0.0063, -0.0054,  0.0193, -0.0018,  0.0124, -0.0141, -0.0186])
```

1.4) Accuracy of the network on the 10000 test images: 45 %

Accuracy per class:

Accuracy for class plane	is: 50.1 %
Accuracy for class car	is: 52.0 %
Accuracy for class bird	is: 46.3 %
Accuracy for class cat	is: 25.4 %
Accuracy for class deer	is: 28.6 %
Accuracy for class dog	is: 42.2 %
Accuracy for class frog	is: 58.9 %
Accuracy for class horse	is: 40.6 %
Accuracy for class ship	is: 46.5 %
Accuracy for class truck	is: 63.9 %

1.5) a)

Accuracy of the network on the 10000 test images: 36 %

Accuracy for class plane	is: 0.0 %
Accuracy for class car	is: 26.0 %
Accuracy for class bird	is: 0.0 %
Accuracy for class cat	is: 25.2 %
Accuracy for class deer	is: 0.0 %
Accuracy for class dog	is: 0.0 %
Accuracy for class frog	is: 22.8 %
Accuracy for class horse	is: 25.6 %
Accuracy for class ship	is: 0.0 %
Accuracy for class truck	is: 0.0 %

b)

Accuracy of the network on the 10000 test images: 20 %

Accuracy for class plane	is: 0.0 %
Accuracy for class car	is: 26.0 %
Accuracy for class bird	is: 0.0 %
Accuracy for class cat	is: 25.2 %
Accuracy for class deer	is: 0.0 %
Accuracy for class dog	is: 0.0 %
Accuracy for class frog	is: 22.8 %
Accuracy for class horse	is: 25.6 %
Accuracy for class ship	is: 0.0 %
Accuracy for class truck	is: 0.0 %


2)

### 2.1) FamNet:

This network architecture has been developed to learn from exemplary in the image itself and count them. It consists of two modules - 1) Feature Extraction 2) Density Map Prediction. Test-time adaption is an important feature where it can count novel classes based on the exemplaries provided in the image. It uses two loss functions 1)Min-count loss 2) Perturbation Loss and combines them to form the Adaptation Loss. We need to tune hyperparameters to make sure combined loss is similar to the training loss so that the gradients get updated properly.

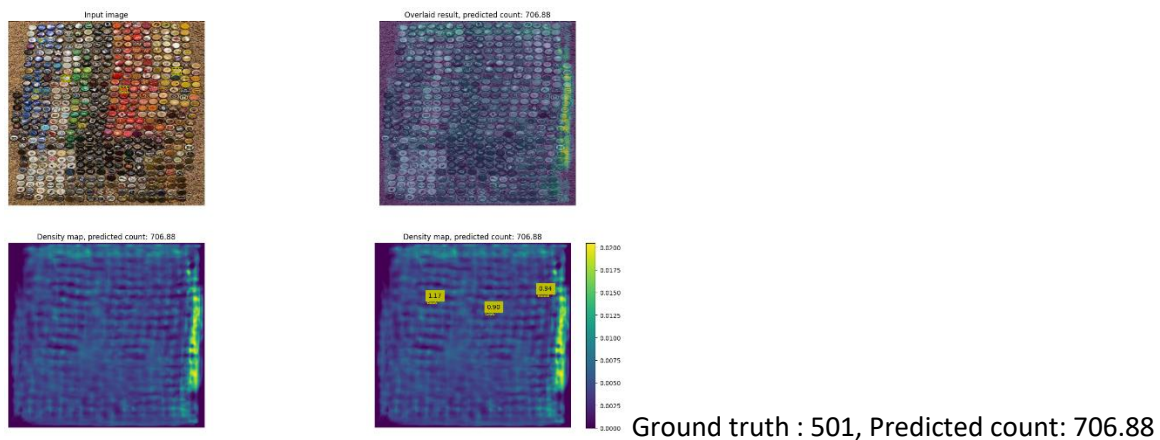
2.2)




Ground truth on X- axis and Predicted count on Y-axis for whole Val Set





On val data, MAE: 24.32, RMSE: 70.94

### Over Count images:










Ground truth : 313, Predicted count:407.36



Ground truth:885 , Predicted Count: 1142.74







Ground truth:501 , Predicted Count: 657.34





Ground truth:60 , Predicted Count: 182.72



### Under-Count images:



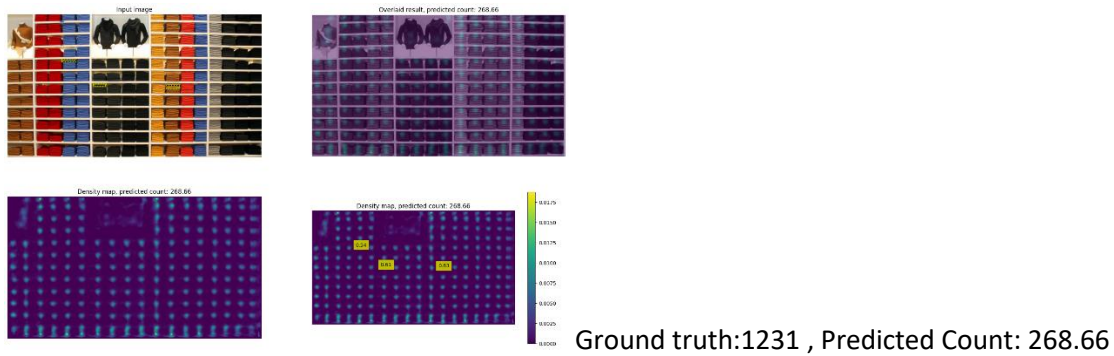
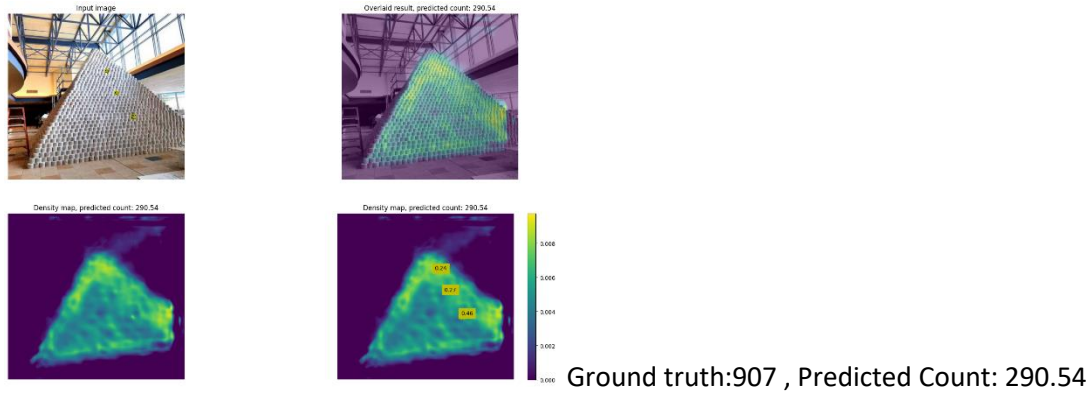
Ground truth:1022 , Predicted Count: 193.42



Ground truth:2092 , Predicted Count: 1148.38



Ground truth:1092 , Predicted Count: 129.48




2.3) With Adaptation,

On val\_PartA data, MAE: 15.70, RMSE: 25.07


Scatter Plot – GT on X-axis, PredCount on Y-axis

1) With Test-time adaptation:





## 2) Without Test-time Adaptation:



Attached `utils.py` and `test.py` since I modified those files only. Changed the `MinLossCount` function to adapt to binary masks information.